


a Four of five scores tie

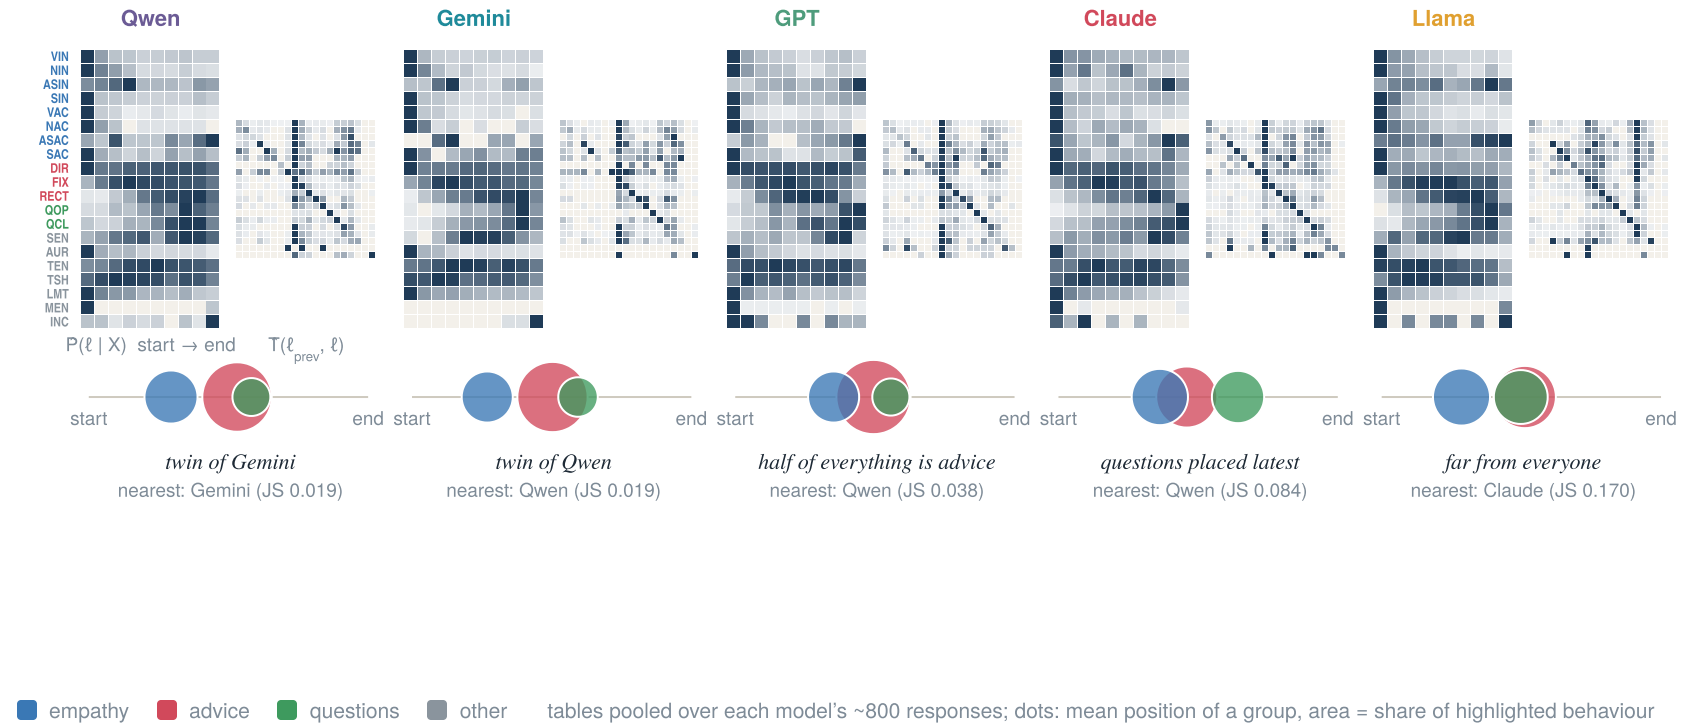
SCORE



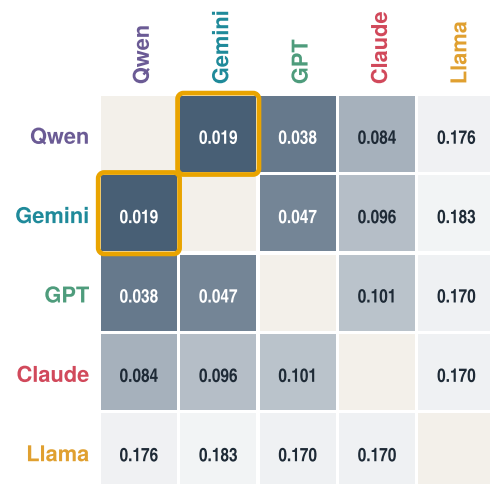
Pooled clinician layer, bootstrap 95% intervals. Four scores overlap (unadjusted $p \geq 0.28$); only Claude separates (Holm $p = 0.025$). The score says **how** scripted, not **whose** script.

b Five script profiles, one order

PROFILE

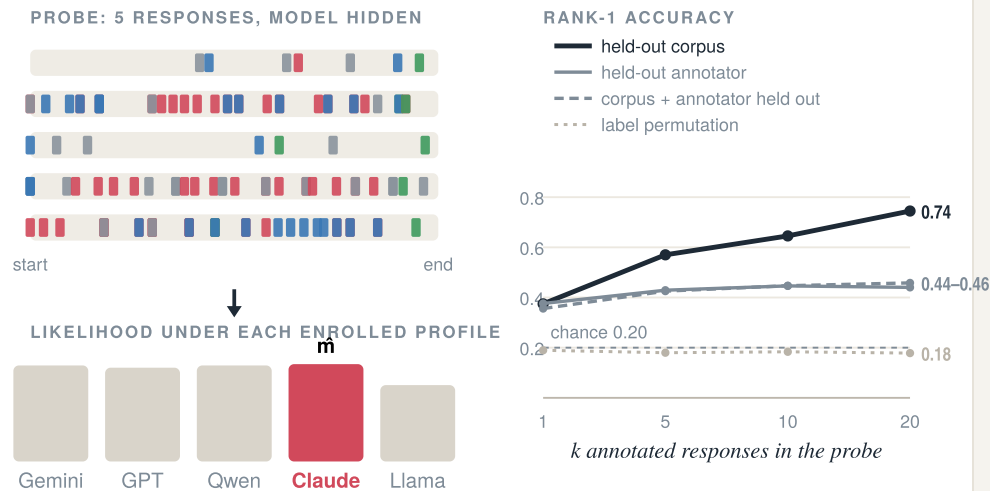


c Profile distance



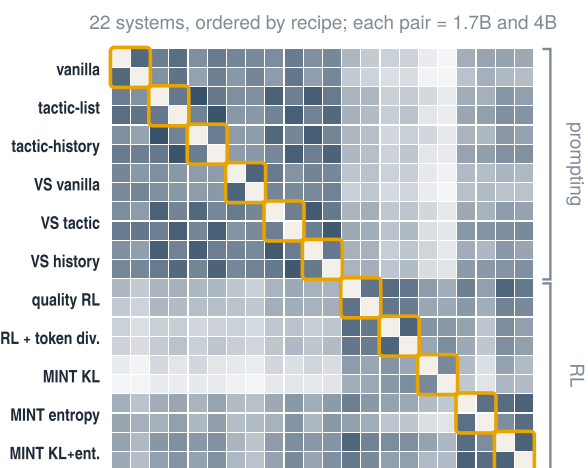
Jensen–Shannon distance between profiles (dark = close). **Qwen–Gemini** sit at 0.019, half the next-closest pair (0.038), and are the closest pair in every one of the four corpora; Llama is ≥ 0.170 from everyone.

d Naming the model from k responses



Left: one real round. Five held-out Claude responses, scored by mean per-event log-likelihood under each profile enrolled on the other three corpora; the highest bar names the model. Right: over 4,000 rounds per point. On a held-out corpus the profile names the model from 20 responses in 74% of rounds; every miss is a Qwen–Gemini confusion.

e Testbed 2: 22 close relatives



Same recipe at the other size (**gold blocks**): mean JS **0.063**. Different recipe at the same size: **0.148**. A 40-turn probe picks its own system in **21/22** (chance 1/22), its own recipe in 21/22, its own size in 22/22: the profile tracks the training objective, not the parameter count.